

Identifying Poisonous Plants Using Deep Learning

Qurat Ul Ain Syeda
MSIS 2626: Deep Learning
Fall 2021



Agenda

01

Defining the Problem

The motivation behind choosing this topic

Data Collection

Using Plant Clef Dataset & Google Images

02

03

Architectures

Simple CNN, Mini-Xception, Transfer Learning & HyperParameter-tuning with KerasTuner

Conclusion

Comparison of results & final model selection

04





Motivation



Background

- ❖ Poisonous plants can cause an allergic reaction in 80-90% of adults
- ❖ Poisonous plants (eg. poison ivy, poison oak, poison sumac) release an oil called urushiol which can cause:
 - Rashes, blisters, bumps are non-threatening
 - Oil can cause lung irritation if inhaled
 - Advisable to see a doctor if too close to eyes or widespread




Use Cases

- ❖ Integrate into nature-related mobile application
- ❖ Eg. hiking trails app, bird watching app, flower/plant identifier app
- ❖ Point, click, and identify functionality

Data Collection

Dataset divided into poisonous and non-poisonous plants

- ❖ Poisonous plants: Downloaded approx. 100 images for each type of plant. Total: 270 images
- ❖ Used PlantCLEF 2015 Annotated Dataset: Consists of over 1000 images submitted by the users of the mobile application Pl@ntNet. Total: 270 images



MODEL ARCHITECTURES



Simple CNN

Using multiple
convolutional layers



Mini-Xception

Truncated version of
Xception



Transfer Learning

Using Xception &
ImageNet

Simple CNN - Workflow

01 **Augment Images w/ ImageDataGenerator()**

Use `.flow_from_directory()` method to load images & apply image augmentation during training

02 **Build, compile & train the model**

Used multiple convolutional layers & RMSProp as an optimization algorithm.

03 **Evaluate model performance & use to predict on test set**

Goal is to maximize validation accuracy



Simple CNN - Model

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 198, 198, 16)	448
max_pooling2d (MaxPooling2D)	(None, 99, 99, 16)	0
conv2d_1 (Conv2D)	(None, 97, 97, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 48, 48, 32)	0
conv2d_2 (Conv2D)	(None, 46, 46, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 23, 23, 64)	0
conv2d_3 (Conv2D)	(None, 21, 21, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 10, 10, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten_1 (Flatten)	(None, 1024)	0
dense_2 (Dense)	(None, 512)	524800
dense_3 (Dense)	(None, 1)	513

Total params: 622,753


Trainable params: 622,753

Non-trainable params: 0




Simple CNN - Training Results

Validation accuracy greater than training due to size & possible difference in noise/variance in validation dataset




Training Loss = 0.485
Validation Loss = 0.413



Training Accuracy = 77.35%
Validation Accuracy = 83.33%


Simple CNN - Predicting on Test Set

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.




Poison Ivy

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.




Poison Oak

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.



Poison Sumac

This image is 100.00 percent non-poisonous and 0.00 percent poisonous.



**Non-Poisonous
Plant**



Simple CNN - Hyperparameter Tuning

- ❖ KerasTuner is an easy-to-use, scalable hyperparameter optimization framework that solves the pain points of hyperparameter search
- ❖ We used the **RandomSearch()** algorithm to find the best hyperparameter values for our model with the goal of maximizing “validation accuracy”
- ❖ The hyperparameters we tuned:
 - # of neurons in the last hidden layer
 - Learning rate supplied to our RMSProp optimizer

```
tuner.search_space_summary()
```

```
Search space summary
```


```
Default search space size: 2
```

```
units (Int)
```

```
{'default': None, 'conditions': [], 'min_value': 32, 'max_value': 512, 'step': 128, 'sampling': None}
```

```
learning_rate (Choice)
```

```
{'default': 0.01, 'conditions': [], 'values': [0.01, 0.001, 0.0001], 'ordered': True}
```



Hyperparameter Tuning Results

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 198, 198, 16)	448
max_pooling2d (MaxPooling2D)	(None, 99, 99, 16)	0
conv2d_1 (Conv2D)	(None, 97, 97, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 48, 48, 32)	0
conv2d_2 (Conv2D)	(None, 46, 46, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 23, 23, 64)	0
conv2d_3 (Conv2D)	(None, 21, 21, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 10, 10, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 160)	164000
dense_1 (Dense)	(None, 1)	161


Total params: 261,601
Trainable params: 261,601
Non-trainable params: 0

Trial summary
Hyperparameters:
units: 160
learning_rate: 0.001
Score: 0.9111110965410868



Training Loss = 0.405

Validation Loss = 0.202



Training Accuracy = 81.43%

Validation Accuracy = 93.33%

Mini-Xception Model- Workflow

01 Load & split dataset using `image_dataset_from_directory()` & augment images using Keras' data augmentation layers

Performed an 80/20 split on dataset & used RandomFlip & RandomRotation layers

02 Build, compile & train the model

Uses pointwise convolution followed by a depthwise convolution

03 Evaluate model performance & use to predict on test set

Goal is to maximize validation accuracy



Xception - Image Augmentation

- ❖ Since we have a limited dataset, we used data augmentation layers.
- ❖ These layers apply random augmentation transforms to a batch of images.
- ❖ They are only active during training.
- ❖ We applied the following transformations:
 - `tf.keras.layers.RandomFlip("horizontal_and_vertical")`
 - `tf.keras.layers.RandomRotation(0.2),`




Xception - Model Summary (Partial)

Model: "model"


Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 180, 180, 3)]	0	[]
sequential (Sequential)	(None, 180, 180, 3)	0	['input_1[0][0]']
rescaling (Rescaling)	(None, 180, 180, 3)	0	['sequential[0][0]']
conv2d (Conv2D)	(None, 90, 90, 32)	896	['rescaling[0][0]']
batch_normalization (Batch Normalization)	(None, 90, 90, 32)	128	['conv2d[0][0]']
activation (Activation)	(None, 90, 90, 32)	0	['batch_normalization[0][0]']
conv2d_1 (Conv2D)	(None, 90, 90, 64)	18496	['activation[0][0]']
batch_normalization_10 (Batch Normalization)	(None, 6, 6, 1024)	4096	['separable_conv2d_8[0][0]']
activation_10 (Activation)	(None, 6, 6, 1024)	0	['batch_normalization_10[0][0]']
global_average_pooling2d (Global Average Pooling2D)	(None, 1024)	0	['activation_10[0][0]']
dropout (Dropout)	(None, 1024)	0	['global_average_pooling2d[0][0]']
dense (Dense)	(None, 1)	1025	['dropout[0][0]']

Total params: 2,782,649
Trainable params: 2,773,913
Non-trainable params: 8,736

Xception Model - Training Results



Training Loss = 0.114
Validation Loss = 0.803




Training Accuracy = 95.65%
Validation Accuracy = 75.23%

Validation accuracy is significantly lower than training due to increased model complexity & limited dataset


Mini-Xception - Predicting on Test Set

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.




Poison Ivy

This image is 0.71 percent non-poisonous and 99.29 percent poisonous.




Poison Oak

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.



Poison Sumac

This image is 99.99 percent non-poisonous and 0.01 percent poisonous.



**Non-Poisonous
Plant**

Transfer Learning - Workflow

01 Load & split dataset using `image_dataset_from_directory()` & augment images using Keras' data augmentation layers

Performed an 80/20 split on dataset & used RandomFlip & RandomRotation layers

02 Build, compile & train the model

- a) Instantiate a base model and load pre-trained weights into it.
- b) Freeze all layers in the base model by setting trainable = False.
- c) Create a new model on top of the output of one (or several) layers from the base model.

03 Evaluate model performance & use to predict on test set

Goal is to maximize validation accuracy




Transfer Learning


Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 180, 180, 3)]	0
sequential (Sequential)	(None, 180, 180, 3)	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
xception (Functional)	(None, 6, 6, 2048)	20861480
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0
dropout (Dropout)	(None, 2048)	0
dense (Dense)	(None, 1)	2049
Total params: 20,863,529		
Trainable params: 2,049		
Non-trainable params: 20,861,480		



Transfer Learning - Training Results



Training Loss = 0.202
Validation Loss = 0.304



Training Accuracy = 92.45%
Validation Accuracy = 87.16%

Validation accuracy is slightly lower than training accuracy

Transfer Learning (Post-Tuning) - Model Summary


Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 180, 180, 3)]	0
sequential (Sequential)	(None, 180, 180, 3)	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
xception (Functional)	(None, 6, 6, 2048)	20861480
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0
dropout (Dropout)	(None, 2048)	0
dense (Dense)	(None, 1)	2049
=====		
Total params: 20,863,529	# of trainable params increases after	
Trainable params: 20,809,001	we unfreeze the model	
Non-trainable params: 54,528		

Transfer Learning (Post-Tuning)

Training Results

Training Accuracy & Validation Accuracy improved slightly




Training Loss = 0.009
Validation Loss = 0.384

Training Accuracy = 99.77%
Validation Accuracy = 90.83%


Transfer Learning- Predicting on Test Set

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.




Poison Ivy

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.




Poison Oak

This image is 0.00 percent non-poisonous and 100.00 percent poisonous.



Poison Sumac

This image is 100.00 percent non-poisonous and 0.00 percent poisonous.





**Non-Poisonous
Plant**

Models Comparison

	Training Accuracy	Validation Accuracy	Training Loss	Validation Loss
Simple CNN	77.35%	83.33%	0.485	0.413
Simple CNN (post-tuning)	81.43%	93.33%	0.405	0.202
Xception	95.65%	75.82%	0.114	0.803
Transfer Learning	92.45%	87.16%	0.202	0.304
Transfer Learning (post-tuning)	99.77%	90.83%	0.009	0.384



Key Takeaways

- ❖ Images can be corrupted during the download process; learnt how to identify & filter them out
 - ❖ Identified various ways to apply data augmentation to enhance dataset & their impact on model performance
 - ❖ Used KerasTuner() to streamline the hyperparameter tuning process
 - ❖ Further enhance the model, using both images and text data to classify various species of plants.
- 
- 

Thanks for Listening!

Any Questions?



References

- **Poisonous Plants:**
 - <https://www.cdc.gov/niosh/topics/plants/default.html>
 - <https://www.webmd.com/allergies/ss/slideshow-poison-plants>
- **Non-poisonous Plants Dataset**
 - <https://www.imageclef.org/lifeclef/2015/plant>
- **Simple CNN:**
 - <https://vijayabhaskar96.medium.com/tutorial-image-classification-with-keras-flow-from-dictionary-and-generators-95f75e5720>
 - <https://www.analyticsvidhya.com/blog/2020/08/image-augmentation-on-the-fly-using-keras-imagedatagenerator/>
- **Hyperparameter Tuning using KerasTuner**
 - https://keras.io/keras_tuner/
- **Mini-Xception Model**
 - https://keras.io/examples/vision/image_classification_from_scratch/#two-options-to-preprocess-the-data
- **Transfer Learning**
 - https://keras.io/guides/transfer_learning/